



كلية هندسة الذكاء الاصطناعي
FACULTY OF ARTIFICIAL INTELLIGENCE ENGINEERING



AI System for Sinus CT Analysis

Graduation Project 2 prepared to obtain the degree of B.Sc. in the Department of Artificial Intelligence Engineering and Data Science

Authors:

Alaa Aldeen Fayez Zaiter

Supervised by:

Dr. Mouhib Alnoukari – Eng. Anas Abdulaziz

١٤٤٧-١٤٤٨ هـ
2025-2026 م • Damascus

SUPERVISOR CERTIFICATION

I certify that the preparation of this project, entitled

AI System for Sinus CT Analysis

prepared by

Alaa Fayez Zaiter

was carried out under my supervision at the Faculty of Artificial Intelligence Engineering in partial fulfillment of the requirements for the degree of B.Sc. in the Department of Artificial Intelligence Engineering and Data Science.

Name:

Date:

Signature:

كلمة شكر وتقدير

الحمد لله رب العالمين، الذي وفقني وأعاني على إتمام هذا العمل، وما كان لهذا الجهد أن يكتمل لولا فضل الله أولاً، ثم دعم وتوجيه أشخاص كان لهم أثر كبير في هذه المسيرة العلمية.

أتقدم بأسمى آيات الشكر والتقدير إلى المهندس أنس عبد العزيز، الذي كان له الدور الأكبر في مرافقتي ودعمي خلال ثلاث سنوات من العمل والتعلم. فقد كان موجهاً صبوراً، وداعماً حقيقياً، لم يخل بعلمه ولا بوقته ولا بنصحه، وكان حاضراً في المراحل الصعبة قبل السهلة. لقد علمني أن النجاح لا يأتي دفعة واحدة، بل يبنى بالصبر، والمثابرة، والإصرار، وأن العمل العلمي يحتاج إلى دقة ومسؤولية وإيمان بالفكرة. فله مني كل الامتنان والتقدير على تحمّله، وتشجيعه، وثقته، وعلى كل ما قدمه من علم وخبرة وتوجيه كان له أثر واضح في إنجاز هذا المشروع.

كما أتقدم بجزيل الشكر والعرفان إلى عميد الكلية الدكتور مهيب النقري، لما قدمه من دعم ورعاية واهتمام، ولما يبذله من جهود في خدمة الكلية وطلابها، وحرصه الدائم على تطوير البيئة العلمية والأكاديمية.

وأتوجه بالشكر والتقدير إلى الدكتور كادان جمعة، والدكتور وسيم الجنيدي نائب عميد الكلية، والدكتور أكرم مسوح، والدكتورة نسرين سليمان، والدكتور رياض سنبل، والمهندسة مروى الصفدي، والمهندس محمد البلخي، لما قدموه من علم وتوجيه ودعم خلال مسيرتي الجامعية، ولما كان لهم من أثر طيب في بناء معرفتي الأكاديمية والعملية.

إلى جميع من وقف إلى جانبي، وساندني بكلمة، أو نصيحة، أو تشجيع، أتقدم بخالص الشكر والامتنان، سائلاً الله أن يجزيهم عني خير الجزاء، وأن يجعل هذا العمل ثمرة طيبة لكل من ساهم في دعمي ومساندتي.

إهداء

إلى والديّ العزيزين، أبي وأمي،

إلى أبي، سندي الأول وظهري الثابت في هذه الحياة، إلى من حمل عني الكثير بصمت، ومنحني القوة والثقة، وكان وجوده مصدر أمانٍ وطمأنينة في كل مرحلة من مراحل طريقي. إلى من علّمني معنى المسؤولية، والصبر، والاعتماد على النفس، وكان دعمه سبباً بعد الله في وصولي إلى هذه اللحظة.

وإلى أُمِّي، نبع الحنان والدعاء، إلى من كان قلبها مأواي، ودعاؤها نوراً يرافقني في كل خطوة، ومنحتني الحب والصبر والطمأنينة دون حدود.

إليكما أهدي هذا العمل، عرفاناً وامتناناً ومحبةً لا تفيها الكلمات.

إلى إخوتي،

شركاء العمر ورفاق الطريق، إلى من كانوا دائماً قريبين من القلب، يشاركونني الفرح، ويخففون عني التعب، ويمنحونني شعور العائلة الذي لا يعوّض.

إلى جدي وجدتي، رمز البركة والحنان، وإلى أعمامي وعماتي، وأخوالي وخالاتي،

بكل ما حملوه لي من محبة ودعم وتشجيع، أهديكم ثمرة هذا الجهد تقديراً لمكانتكم الغالية في قلبي.

إلى أحبتي وأصدقائي وزملائي،

الذين كانوا عوناً في اللحظات الصعبة، ورفاقاً في أيام التعب والإنجاز، إلى من شاركوني هذه الرحلة بكلمة طيبة، أو موقف صادق، أو دعم لا يُنسى وأخص بالشكر أصدقائي الأعزاء وزملائي: م. مؤمن هشام رباح و محمد عمر باهر الشعار، م. أنس عوض، د. احمد حسن، يمان فضل الله، م. معاذ جادالله، م. محمد جدوع، م. عبدالملك سلامة، م. محمد البلخي، م. خالد المصري، م. مضر الصالح، م. محمد الصباغ، م. احمد خلوف، م. حسان الأنقر، م. محمد الحمود، م. رواد الزعبي، م. احمد عمشة، م. كاسم المليح، م. محمد الموسى، م. منار طعمة، م. محمد زاهر كريم، م. آدم العُميان، م. محمد يحيى أبو سمرة، م. عثمان دركل، م. محمد اواب خير، م. فراس الزعبي، م. تقى الزعبي، م. نور الهدى نعتسان، م. فادي الخوري حنا، م. طارق المحيسن، م. مهند المبيض، م. حسن الحمصي، م. عبدالرحمن العامري.

إلى كل من كان له أثر جميل في هذه المسيرة، أهدي هذا العمل بمحبة وامتنان.

الملخص

يعالج هذا المشروع مشكلة تحليل صور التصوير الطبقي المحوري للجيوب الأنفية باستخدام تقنيات التعلم العميق. وتكمن أهمية هذه المسألة في أن صور الجيوب الأنفية تحتوي على بُنى تشريحية متداخلة، كما أن تقييمها يدوياً قد يكون مرهقاً ويتطلب خبرة اختصاصية. وتظهر تحديات إضافية بسبب محدودية البيانات، وعدم توازن الحالات المرضية، واحتمال ترافق أكثر من شذوذ تشريحي أو مرضي في الحالة الواحدة.

يقترح المشروع خط معالجة يعتمد أولاً على التقسيم التشريحي باستخدام نموذج SwinUNETR ، وذلك لتحديد مناطق مهمة مثل الجيب الفكي الأيمن، والجيب الفكي الأيسر، والتجويف الأنفي الأيمن، والتجويف الأنفي الأيسر، والبلعوم الأنفي. بعد ذلك، يتم استخراج مناطق اهتمام موجّهة تشريحياً، ثم تدريب نموذج تصنيف متعدد التسميات للكشف عن خمس علامات مرضية مختارة. وقد تم اعتماد نموذج ConvNeXt-Tiny كنموذج التصنيف النهائي بعد مقارنته مع نموذج Swin-T.

أظهرت نتائج التقسيم قدرة جيدة على تحديد البنى التشريحية، حيث بلغ متوسط معامل Dice النهائي 87.43%، وبلغ متوسط IoU مقدار 81.02% على مجموعة الاختبار. أما في مرحلة التصنيف، فقد حقق نموذج ConvNeXt-Tiny أداءً جيد وفقاً لمتوسط Macro F1 عبر الطيات حيث بلغ 78.49% ونتائج التنبؤ خارج الطية 69.28% OOF F1 . ومع ذلك، بقيت بعض التسميات، وخاصة شذوذ الجيب الفكي، أقل استقراراً بسبب انخفاض عدد العينات الإيجابية وترافقها المتكرر مع شذوذات أخرى.

Abstract

This project addresses the problem of artificial-intelligence-assisted analysis of sinus computed tomography (CT) scans. Sinus CT interpretation is clinically important because it supports the assessment of sinonasal anatomy and pathology, but automatic analysis is challenging due to the three-dimensional nature of CT volumes, anatomical variability, limited annotated data, and strong class imbalance among abnormal findings.

The proposed solution follows a segmentation-assisted pipeline. First, a SwinUNETR model is trained to segment key anatomical structures in NasalSeg CT volumes. Second, anatomy-guided regions of interest (ROIs) are extracted from the CT scan and segmentation masks. Third, a multi-label classifier predicts five final abnormality labels: nasal septum abnormality, right maxillary abnormality, left maxillary abnormality, right inferior turbinate hypertrophy, and left inferior turbinate hypertrophy. ConvNeXt-Tiny is selected as the final classifier after comparison with Swin-T.

The segmentation model achieved a mean Dice score of 0.8743 and mean IoU of 0.8102 on the test set. For classification, ConvNeXt-Tiny achieved a fold-level macro F1 score of 0.7849 and an out-of-fold macro F1 score of 0.6928, outperforming Swin-T in the final comparison. The results show that the proposed pipeline is promising, while also revealing important limitations caused by small positive sample counts, multi-label co-occurrence, and indirect anatomical targets such as turbinate abnormalities.

Table of Contents

Authors:	I
Supervised by:	I
SUPERVISOR CERTIFICATION.....	II
كلمة شكر وتقدير	III
إهداء	IV
الملخص	V
Abstract.....	VI
Table of Contents.....	VII
List of Tables	X
List of Figures.....	XI
List of Equations.....	XII
Glossary of Terms and Symbols	XIII
Chapter One: Introduction.....	15
1.1 Introduction	15
1.2 Scientific Problem of the Research Project	15
1.3 Research Objectives	16
1.4 Research Contribution.....	16
1.5 Target Beneficiaries	17
1.6 Report Organization	17
Chapter Two: Theoretical Background	18
2.1 Computed Tomography and Sinus Imaging.....	18
2.2 Sinonasal Anatomy Relevant to the Project	19
2.3 Medical Image Segmentation	19
2.4 Deep Learning for CT Analysis.....	20
2.5 Convolutional Neural Networks and ConvNeXt.....	20
2.6 Vision Transformers, Swin Transformer, and SwinUNETR	22
2.7 Multi-Label Classification.....	22
2.8 Class Imbalance and Focal Loss	23
2.9 Evaluation Metrics	23
2.10 Chapter Summary	25
Chapter Three: Literature Review.....	26
3.1 Introduction.....	26
3.2 Previous Studies in Sinus CT Segmentation	26
3.3 NasalSeg Dataset and Benchmark	26
3.4 SwinUNETR for 3D Medical Image Segmentation	27
3.5 Whangbo et al. Paranasal Sinus Segmentation.....	27

3.6 Ozturk et al. Maxillary Sinus Segmentation.....	27
3.7 Song et al. CNN, Transformer, and Hybrid Segmentation Comparison.....	27
3.8 Previous Studies in Sinus and Paranasal Abnormality Classification	28
3.9 Bhattacharya et al. Self-Supervised Maxillary Anomaly Classification	28
3.10 Supervised Contrastive Learning for Maxillary Anomalies	28
3.11 Classification Backbones and Transfer Learning Studies	28
3.12 ConvNeXt.....	29
3.13 Swin Transformer	29
3.14 MedNeXt and ConvNeXt-Style Medical Networks	29
3.15 Improved ConvNeXt-Tiny and CT Classification Studies	29
3.16 2.5D Medical Imaging Approaches	30
3.17 State-of-the-Art Frameworks Analysis.....	30
3.18 Comparative Analysis of Existing Systems.....	33
3.19 Research Gap and Proposed Innovation	34
3.20 Chapter Summary	35
Chapter Four: Research Methodology	36
4.1 Introduction	36
4.2 Dataset Description and Annotation Characteristics	36
4.3 Data Preprocessing.....	38
4.4 Data Augmentation	39
4.5 Data Partitioning and 5-Fold Evaluation.....	39
4.6 SwinUNETR Segmentation Stage	40
4.7 ROI Extraction Strategy	41
4.8 2.5D ROI Generation	43
4.9 Important Exploratory Experiments	45
4.10 Final Selected Classification Model	49
4.11 Implementation Environment.....	49
4.12 Chapter Summary and Research Contribution	50
Chapter Five: Research Results.....	51
5.1 Introduction	51
5.2 Tools and Evaluation Methods.....	51
5.3 Segmentation Results	52
5.4 Classification Results	54
5.5 Comparison with Previous Works	57
5.6 Discussion of Results	58
5.7 Chapter Summary	60
Chapter Six: Conclusion.....	61

6.1	Introduction	61
6.2	Conclusions	61
6.3	Research Limitations.....	61
6.4	Challenges and Future Work.....	62
	References	64

List of Tables

Table 0.1: Glossary of terms and abbreviations.....	XIII
Table 3.1: Summary of related work in sinus CT analysis	28
Table 3.2: Comparison of existing approaches and the proposed system	30
Table 4.1: Dataset summary	33
Table 4.2: Final target labels and positive sample counts	33
Table 4.3: Data partitioning strategy used in this project	35
Table 4.4: SwinUNETR segmentation training configuration	35
Table 4.5: ROI-to-label mapping used in the final classification system	36
Table 4.6: Timeline of major exploratory experiments and outcomes	41
Table 5.1: Overall segmentation performance on the test set	45
Table 5.2: Per-class segmentation performance	45
Table 5.3: Supporting ConvNeXt-Tiny and Swin-T comparison	46
Table 5.4: OOF per-label metrics for ConvNeXt-Tiny	47
Table 5.5: Comparison of the proposed model with previous works	50
Table 6.1: Summary of main findings and future work directions	54

List of Figures

Figure 2.1: Main sinonasal anatomical regions in CT imaging	18
Figure 2.2: Example CT views of the sinonasal area, coronal, sagittal, axial	19
Figure 2.3: Simplified ConvNeXt-Tiny classification concept	20
Figure 2.4: Simplified SwinUNETR architecture overview	21
Figure 3.1: Research gap and positioning of the proposed system	27
Figure 3.2: ConvNeXt-Tiny and Swin-T literature justification	27
Figure 4.1: Overall pipeline of the proposed sinus CT analysis system	32
Figure 4.2: Label preparation workflow from the annotation sheet	34
Figure 4.3: Augmentation sanity-check visualization	36
Figure 4.4: Final ROI definitions used in the classification stage	37
Figure 4.5: Example ROI extraction results from the classification notebook	38
Figure 4.6: 2.5D ROI generation from axial, coronal, and sagittal views	38
Figure 4.7: Timeline of major classification experiments	40
Figure 4.8: Macro F1 comparison across experiments	42
Figure 5.1: SwinUNETR training and validation curves	45
Figure 5.2: Example segmentation prediction on a test case	46
Figure 5.3: ConvNeXt-Tiny loss curves	47
Figure 5.4: ConvNeXt-Tiny macro F1 curves	47
Figure 5.5: ConvNeXt-Tiny 5×5 dominant-label confusion matrix	48
Figure 5.6: ROC curves for the final ConvNeXt-Tiny model	49
Figure 5.7: Precision-Recall curves for the final ConvNeXt-Tiny model	49

List of Equations

Equation [2.1]: Dice Similarity Coefficient.....	19
Equation [2.2]: Intersection over Union	20
Equation [2.3]: Binary classification decision from sigmoid probability.....	21
Equation [2.4]: Binary focal loss concept.....	22
Equation [2.5]: Precision	22
Equation [2.6]: Recall / Sensitivity	22
Equation [2.7]: Specificity	22
Equation [2.8]: F1-score	23
Equation [2.9]: Balanced accuracy	23
Equation [2.10]: Exact match accuracy	23

Glossary of Terms and Symbols

Table 0.1: Glossary of terms and abbreviations

Term	English Equivalent / Description	Symbol / Abbreviation
Computed Tomography	Cross-sectional X-ray imaging modality used to visualize internal anatomy.	CT
Region of Interest	A selected image region used as model input for a specific label.	ROI
Dice Similarity Coefficient	Overlap metric used to evaluate segmentation masks.	Dice
Intersection over Union	Overlap metric between predicted and reference masks.	IoU
Convolutional Neural Network	Deep learning architecture based on convolutional operations.	CNN
Vision Transformer	Neural architecture that applies self-attention to image patches.	ViT
SwinUNETR	Transformer-based 3D U-shaped segmentation architecture.	SwinUNETR
ConvNeXt	Modern convolutional architecture inspired by transformer-era design choices.	ConvNeXt
Out-of-Fold Prediction	Prediction produced for each sample by a model that did not train on it.	OOF
Focal Loss	Classification loss that reduces the contribution of easy examples.	Focal Loss
Average Precision	Area under the precision-recall curve, useful for imbalanced classification.	AP
Exact Match Accuracy	Multi-label metric requiring all labels of a case to be predicted correctly.	Subset Accuracy

Chapter One: Introduction

1.1 Introduction

Computed tomography (CT) is one of the most important imaging modalities for evaluating the paranasal sinuses and nasal cavity. It provides detailed cross-sectional views of bony and soft-tissue structures, making it useful for diagnosing sinus disease, studying anatomical variation, and planning endoscopic sinus surgery. However, CT volumes are three-dimensional and usually contain many slices, which makes manual review time-consuming and dependent on clinical experience.

Artificial intelligence, especially deep learning, has become an important research direction in medical image analysis. Deep learning models can learn visual patterns directly from imaging data and can be trained to perform segmentation, classification, detection, and quantification tasks. In sinus CT analysis, segmentation can localize anatomical structures, while classification can support the detection of abnormal findings. Combining these two tasks may help produce a system that is more anatomically meaningful than a purely image-level classifier.

This project focuses on developing an AI system for sinus CT analysis using a segmentation-assisted classification pipeline. The system first localizes important sinonasal structures using a SwinUNETR segmentation model, then extracts anatomy-guided regions of interest, and finally classifies five selected abnormality labels using modern image-classification backbones. The final classification model selected in this work is ConvNeXt-Tiny because it achieved the strongest performance among the final evaluated classifiers.

1.2 Scientific Problem of the Research Project

The scientific problem addressed in this project is the automatic analysis of sinus CT scans for selected anatomical and pathological abnormalities. The problem is not a simple image classification task because the abnormalities are region-dependent, multi-label, and often correlated. A patient may have septal deviation, turbinate hypertrophy, and maxillary sinus abnormality at the same time. A single-label classifier or a simple dominant-label confusion matrix cannot fully represent this situation.

The problem also includes data-related challenges. The dataset contains a limited number of CT

cases, and the positive counts differ strongly across labels. For example, the final labels include 71 positive cases for nasal septum abnormality, 60 positive cases for each inferior turbinate hypertrophy label, but only 22 and 23 positives for the right and left maxillary abnormality labels. This imbalance makes maxillary abnormalities less stable during training and evaluation.

Another difficulty is the mismatch between available segmentation masks and some clinical targets. NasalSeg provides masks for the right and left maxillary sinuses, right and left nasal cavities, and nasopharynx. It does not directly segment the nasal septum or turbinates. Therefore, the final model must infer septal and turbinate-related abnormalities using contextual neighboring regions rather than direct target masks.

1.3 Research Objectives

- Develop an AI-assisted pipeline for sinus CT analysis based on segmentation-guided region extraction.
- Train a 3D SwinUNETR model to segment key nasal and paranasal anatomical structures.
- Convert clinical annotation-sheet findings into binary abnormality labels suitable for multi-label learning.
- Design stable anatomical ROIs that connect segmentation outputs to clinically meaningful classification targets.
- Compare ConvNeXt-Tiny and Swin-T as final classification backbones under a 5-fold evaluation setup.
- Select the best-performing final classifier and analyze its strengths, weaknesses, and failure cases.

1.4 Research Contribution

The main contribution of this project is the development of a practical segmentation-assisted multi-label classification workflow for sinus CT analysis. Instead of classifying the full CT volume blindly, the model uses anatomical localization to extract targeted ROIs. This helps connect each classification output to a relevant anatomical area and makes the pipeline easier to interpret.

A second contribution is the final ROI design. Because NasalSeg does not directly provide septum or turbinate masks, the project explores contextual regions based on the available nasal cavity masks. The final septum ROI is defined using both nasal cavities with padding, representing the idea that the septum lies between the right and left nasal cavity regions. The inferior turbinate labels are modeled using the ipsilateral nasal cavity ROIs.

A third contribution is the comparative evaluation of ConvNeXt-Tiny and Swin-T for this specific multi-label sinus CT task. ConvNeXt-Tiny is selected as the final model because it achieves better fold-level and out-of-fold performance than Swin-T. The report also discusses why some labels, especially maxillary abnormalities, remain unstable due to low positive counts and label co-occurrence.

1.5 Target Beneficiaries

- Radiologists and ENT physicians who interpret sinus CT scans.
- Hospitals and imaging centers interested in AI-assisted triage or decision support.
- Researchers working on medical image segmentation and classification.
- Students and engineers studying applied deep learning in medical imaging.
- Developers who may extend the system into a clinical review or educational tool.

1.6 Report Organization

Chapter Two presents the theoretical background needed to understand the project, including CT imaging, sinus anatomy, segmentation, classification, CNNs, transformers, and evaluation metrics. Chapter Three reviews related work and identifies the research gap. Chapter Four explains the research methodology, dataset, preprocessing, segmentation model, ROI extraction, and final classification models. Chapter Five presents the experimental results and discusses the segmentation and classification outcomes. Chapter Six concludes the report and proposes future work.

Chapter Two: Theoretical Background

2.1 Computed Tomography and Sinus Imaging

Computed tomography constructs cross-sectional images from X-ray measurements acquired around the patient. In sinus imaging, CT is especially valuable because it shows the bony boundaries of the paranasal sinuses and nasal cavity. The intensity values in CT images are commonly represented in Hounsfield Units (HU), where air has very low values and bone has high values. This contrast is useful for distinguishing air-filled cavities from surrounding structures.

Sinus CT volumes are usually analyzed in axial, coronal, and sagittal planes. A model trained on CT data must therefore learn three-dimensional spatial relationships, or approximate them through multi-view two-dimensional representations. In this project, segmentation is performed as a 3D task, while classification uses 2.5D ROI images derived from each anatomical crop.




Figure 2.1: Main sinonasal anatomical regions in CT imaging




Figure 2.2: Example CT views of the sinonasal area, coronal, sagittal, axial

2.2 Sinonasal Anatomy Relevant to the Project

The paranasal sinuses are air-filled cavities surrounding the nasal cavity. The maxillary sinuses are located on both sides of the nasal cavity and are among the largest paranasal sinuses. The nasal cavity is divided into right and left sides by the nasal septum. The turbinates are curved bony structures covered by mucosa that project into the nasal cavity and affect airflow.

The final classification labels selected in this project are connected to these anatomical regions. Right and left maxillary abnormalities are associated with the corresponding maxillary sinus ROIs. Nasal septum abnormality is associated with the region between the two nasal cavities. Inferior turbinate hypertrophy is associated with the right or left nasal cavity because direct turbinate segmentation masks are not available in the dataset.

2.3 Medical Image Segmentation

Medical image segmentation is the task of assigning a class label to each pixel or voxel. In 3D CT segmentation, each voxel in the volume is assigned to an anatomical class such as background, right maxillary sinus, left maxillary sinus, right nasal cavity, left nasal cavity, or nasopharynx. Segmentation is useful because it localizes structures and enables measurement, visualization, and region-specific analysis.

In this project, segmentation is not the final clinical output. Instead, it is an intermediate stage used to guide ROI extraction. This design allows the classification model to focus on relevant

anatomical regions rather than the entire CT volume.

$$Dice = 2 |P \cap G| / (|P| + |G|)$$

Equation [2.1]: Dice Similarity Coefficient

$$IoU = |P \cap G| / |P \cup G|$$

Equation [2.2]: Intersection over Union

2.4 Deep Learning for CT Analysis

Deep learning models learn hierarchical representations from data. Early layers usually capture simple intensity and edge patterns, while deeper layers capture more complex anatomical or pathological features. For CT analysis, the model must be robust to anatomical variation, acquisition differences, and class imbalance.

There are two broad approaches to CT learning. A 3D model processes the full volume or volume crops and can preserve volumetric context, but it requires more memory and more data. A 2D or 2.5D model processes slices or multi-plane views, reducing memory requirements and enabling the use of pretrained natural-image backbones. This project uses 3D modeling for segmentation and 2.5D transfer learning for classification.

2.5 Convolutional Neural Networks and ConvNeXt

Convolutional neural networks are well suited for medical images because they exploit local spatial structure. ConvNeXt is a modern convolutional architecture that updates classical CNN design choices while preserving the simplicity and locality of convolutional operations. In this project, ConvNeXt-Tiny is used as the final classification backbone because it showed the best balance between performance and stability on the selected multi-label task.

The selection of ConvNeXt-Tiny is supported by two ideas. First, ConvNeXt preserves convolutional inductive bias, which is often helpful in limited-data medical imaging. Second, its design includes transformer-era improvements while avoiding some of the data demands of pure

transformer backbones. This makes it a reasonable candidate for ROI-based 2.5D classification, where the input already focuses on a clinically meaningful region.




Figure 2.3: Simplified ConvNeXt-Tiny classification concept

2.6 Vision Transformers, Swin Transformer, and SwinUNETR

Vision transformers process images through patch representations and self-attention. Standard transformers can be computationally expensive for high-resolution medical images. Swin Transformer addresses this issue by using hierarchical shifted windows, which makes attention more efficient while still capturing contextual information. SwinUNETR combines a Swin Transformer encoder with a U-shaped segmentation decoder, making it suitable for 3D medical segmentation tasks.




Figure 2.4: Simplified SwinUNETR architecture overview

2.7 Multi-Label Classification

In multi-label classification, each case can belong to more than one class. This is different from multi-class classification, where each sample belongs to exactly one class. The sinus CT task in this project is multi-label because one case may contain multiple simultaneous abnormalities. For this reason, each final label is predicted with an independent sigmoid output rather than a single softmax output.

$$\hat{y}_i = 1 \text{ if } \text{sigmoid}(z_i) \geq \tau_i, \text{ otherwise } 0$$

Equation [2.3]: Binary classification decision from sigmoid probability

2.8 Class Imbalance and Focal Loss

Class imbalance occurs when positive and negative examples are unevenly distributed. In the final dataset, maxillary abnormality labels have much fewer positive samples than septum and inferior turbinate labels. This can cause the model to overfit rare labels or predict them inconsistently. Focal loss helps by reducing the effect of easy examples and giving more attention to difficult cases. The final experiments use focal loss with gamma equal to 1 and positive class weighting.

$$FL(p_t) = -\alpha_t (1 - p_t)^\gamma \log(p_t)$$

Equation [2.4]: Binary focal loss concept

2.9 Evaluation Metrics

2.9.1 Segmentation Metrics

The segmentation model was evaluated with spatial overlap metrics and runtime measures. Dice and IoU measure agreement between predicted and reference masks. Precision measures the proportion of predicted foreground voxels that are correct, while recall or sensitivity measures the proportion of reference voxels captured by the model. Inference time and throughput describe the practical speed of the trained model during test-time prediction.

$$Precision = TP / (TP + FP)$$

Equation [2.5]: Precision

$$Recall = TP / (TP + FN)$$

Equation [2.6]: Recall / Sensitivity

$$\textit{Specificity} = TN / (TN + FP)$$

Equation [2.7]: Specificity

2.9.2 Classification Metrics

The classification task was evaluated using complementary metrics because the task is imbalanced and multi-label. Accuracy measures label-wise correctness. Precision and recall measure false-positive and false-negative behavior. Specificity measures the ability to reject negative cases. Balanced accuracy averages sensitivity and specificity, making it more informative when positive and negative cases are imbalanced.

F1-score is the harmonic mean of precision and recall. Macro F1 averages label-wise F1 scores, giving each abnormality equal importance regardless of its frequency. AUC measures ranking quality across thresholds, while average precision summarizes the precision-recall curve and is useful for imbalanced labels. Overall label accuracy measures correctness over all labels, while exact match accuracy requires every label of a case to be predicted correctly.

The report also uses out-of-fold evaluation. In a 5-fold setting, each case is predicted by a model that did not train on it. This produces a more reliable estimate of generalization than evaluating

only on the training folds. Threshold tuning is performed on the OOF predictions to select label-specific probability cutoffs. The 5×5 dominant-label confusion matrix is used only as a visualization because it compresses a multi-label prediction into one dominant true label and one dominant predicted label.

$$F1 = 2 \times Precision \times Recall / (Precision + Recall)$$

Equation [2.8]: F1-score

$$Balanced\ Accuracy = (Sensitivity + Specificity) / 2$$

Equation [2.9]: Balanced accuracy

$$Exact\ Match = \text{number of cases with all labels correct} / \text{total number of cases}$$

Equation [2.10]: Exact match accuracy

2.10 Chapter Summary

This chapter introduced the theoretical foundations of the project, including sinus CT, sinonasal anatomy, segmentation, multi-label classification, CNNs, transformers, class imbalance, and evaluation metrics. These concepts support the methodology described in Chapter Four.

Chapter Three: Literature Review

3.1 Introduction

This chapter reviews previous studies related to sinus CT segmentation, paranasal anomaly classification, medical-image backbones, and 2.5D transfer learning. Similar to the organization used in the ChromoSep-AI literature review, the chapter discusses the most relevant studies separately rather than reducing them to a single table. This structure makes it clearer how each work influenced the design of the proposed segmentation-assisted sinus CT analysis system.

The reviewed literature shows that the field contains strong segmentation studies and strong general classification backbones, but fewer works combine anatomical sinus segmentation with multi-label abnormality prediction. Therefore, the review focuses on the research gap between segmentation-only systems and abnormality-classification systems.

3.2 Previous Studies in Sinus CT Segmentation

Previous sinus CT segmentation studies mainly focused on localizing anatomical structures such as the maxillary sinuses, nasal cavity, nasopharynx, and other paranasal sinus regions. These works are important for this project because anatomical localization is the first stage of the proposed pipeline. Instead of using segmentation as the final clinical output, this project uses segmentation to guide ROI extraction before abnormality classification. The reviewed segmentation papers therefore provide the technical foundation for using 3D models such as SwinUNETR and U-Net variants in sinonasal CT analysis.

3.3 NasalSeg Dataset and Benchmark

Zhang et al. introduced NasalSeg, an open-access dataset for nasal cavity and paranasal sinus segmentation from 3D CT images [1]. The dataset contains 130 CT scans with voxel-wise annotation of five structures: right maxillary sinus, left maxillary sinus, right nasal cavity, left nasal cavity, and nasopharynx. This work is the most directly related foundation for the present project because the proposed system uses the same anatomical structures to guide ROI extraction. However, NasalSeg does not directly annotate the nasal septum or inferior turbinates, which creates an important limitation for downstream abnormality classification.

3.4 SwinUNETR for 3D Medical Image Segmentation

Hatamizadeh et al. proposed SwinUNETR, a transformer-based encoder-decoder architecture for 3D medical image segmentation [3]. The model uses a Swin Transformer encoder to capture hierarchical volumetric context and a U-shaped decoder to reconstruct dense segmentation masks. In their medical segmentation benchmarks, transformer-based 3D modeling showed strong Dice performance and motivated its use for volumetric anatomical localization. In this project, SwinUNETR was selected for the segmentation stage because sinus CT volumes require spatial context across slices.

3.5 Whangbo et al. Paranasal Sinus Segmentation

Whangbo et al. compared several 3D U-Net variants for multi-class paranasal sinus segmentation in sinusitis CT images [6]. The normal 3D U-Net achieved the best reported F1 score, reaching 84.29% on the normal test set and 79.32% on the abnormal test set. This study supports the use of 3D CNN-based segmentation for sinus anatomy, especially when sinusitis changes the appearance of the target regions. Its limitation for the present project is that it focuses on anatomical segmentation rather than multi-label abnormality prediction.

3.6 Ozturk et al. Maxillary Sinus Segmentation

Ozturk et al. developed a U-Net-based method for automatic maxillary sinus segmentation on cone-beam CT images [9]. The study is relevant because maxillary sinus abnormality is one of the final target groups in this project. It demonstrates that deep learning can localize maxillary sinus boundaries effectively, which is important for region-specific analysis. However, the work is CBCT-focused and does not address a full sinus CT pipeline with septum, turbinate, and bilateral multi-label abnormalities.

3.7 Song et al. CNN, Transformer, and Hybrid Segmentation Comparison

Song et al. compared CNN, vision-transformer, and hybrid networks for paranasal sinus segmentation in CT images [10]. Their results showed that Swin UNETR achieved the strongest overall segmentation performance among the tested models, with a mean Jaccard Index of 0.719, Dice Similarity Coefficient of 0.830, recall of 0.758, and HD95 of 10.529. This study is important because it directly compares architecture families in the sinus domain. It also supports the idea that hybrid 3D architectures are strong candidates for sinus segmentation, although it

remains focused on segmentation rather than diagnosis.

3.8 Previous Studies in Sinus and Paranasal Abnormality Classification

Previous classification studies in sinus and paranasal imaging have commonly focused on narrower diagnostic targets, especially maxillary sinus anomaly detection. These studies show that deep learning can support abnormality recognition in CT or CBCT images, but most of them do not address a five-label multi-label setting. In contrast, the current project predicts septum abnormality, bilateral maxillary abnormality, and bilateral inferior turbinate hypertrophy from the same CT case. This makes the classification problem more difficult because labels may co-occur and some target structures are not directly segmented.

3.9 Bhattacharya et al. Self-Supervised Maxillary Anomaly Classification

Bhattacharya et al. proposed a self-supervised learning approach for classifying paranasal anomalies in the maxillary sinus [7]. The method used a 3D convolutional autoencoder and a localization-oriented self-supervised task before fine-tuning for normal-versus-anomalous maxillary sinus classification. When trained with only 10% of the annotated dataset, the proposed method achieved an AUPRC of 0.79, outperforming BYOL, SimSiam, SimCLR, and SparK baselines. This is relevant to the present project because it addresses limited annotation, but it remains narrower because it focuses on maxillary sinus anomalies rather than five-label sinus CT classification.

3.10 Supervised Contrastive Learning for Maxillary Anomalies

Bhattacharya et al. also studied supervised contrastive learning for classifying paranasal anomalies in the maxillary sinus [8]. Their 3D CNN trained with supervised contrastive learning and cross-entropy achieved an AUROC of 0.85, compared with 0.66 for a 3D CNN trained with cross-entropy alone. This result shows that representation learning can improve classification in limited-data paranasal anomaly settings. However, the task was still focused on maxillary sinus anomaly classification and did not solve the broader problem of multi-label sinus CT abnormality prediction.

3.11 Classification Backbones and Transfer Learning Studies

The classification stage of this project is also connected to broader backbone and transfer-learning studies. Modern CNNs such as ConvNeXt and hierarchical transformers such as Swin

Transformer provide strong pretrained feature extractors that can be adapted to medical images. Because the available dataset is small, transfer learning is more practical than training a large classifier from scratch. The reviewed backbone studies justify the final comparison between ConvNeXt-Tiny and Swin-T and support the use of 2.5D ROI images as inputs to pretrained 2D models.

3.12 ConvNeXt

Liu et al. introduced ConvNeXt as a modernized convolutional architecture that adopts several transformer-era design choices while preserving the inductive bias of CNNs [4]. The ConvNeXt family achieved competitive large-scale vision performance, with the overall family reaching 87.8% ImageNet top-1 accuracy in the original work. This is relevant to the present project because the final classifier uses localized 2.5D ROI images and limited medical data, where convolutional locality can be beneficial. ConvNeXt-Tiny was therefore selected as the final backbone after outperforming the transformer comparison in the project experiments.

3.13 Swin Transformer

Liu et al. proposed Swin Transformer, a hierarchical vision transformer based on shifted windows [5]. The model achieved 87.3% ImageNet-1K top-1 accuracy and strong dense-prediction results, including 53.5 mIoU on ADE20K validation in the original paper. Swin is relevant because it represents the transformer family and motivates SwinUNETR for segmentation and Swin-T as a classification comparison. In this project, Swin-T is treated as an important comparison experiment rather than the final selected classifier.

3.14 MedNeXt and ConvNeXt-Style Medical Networks

Roy et al. introduced MedNeXt, a ConvNeXt-inspired architecture for medical image segmentation [12]. The model adapts large-kernel convolutional design to 3D encoder-decoder segmentation and shows that modern CNN principles can remain competitive in medical imaging. This supports the broader idea that convolutional models should not be dismissed in favor of transformers only. For the present project, this is conceptually important because ConvNeXt-Tiny achieved the best classification behavior on the limited ROI-based dataset.

3.15 Improved ConvNeXt-Tiny and CT Classification Studies

Recent medical classification studies have also adapted ConvNeXt-Tiny for efficient medical-

image recognition [11]. FocalNeXt further shows that ConvNeXt-related components can be integrated into CT classification systems [13]. These studies are not sinus-specific, but they support the use of lightweight modern convolutional backbones in limited medical datasets. This literature is consistent with the final project result, where ConvNeXt-Tiny gave better fold-level and out-of-fold performance than Swin-T.

3.16 2.5D Medical Imaging Approaches

Several medical-imaging studies use 2.5D representations to balance volumetric context and computational efficiency [15]. In this approach, representative slices or orthogonal views are used as multi-channel inputs for 2D pretrained networks. This strategy is relevant to the current project because each 3D ROI is converted into axial, coronal, and sagittal views. The resulting representation preserves partial 3D context while enabling transfer learning with pretrained ConvNeXt-Tiny and Swin-T backbones.

3.17 State-of-the-Art Frameworks Analysis

The current state of the art in sinus CT analysis can be understood as two related directions. The first direction is high-quality anatomical segmentation, where 3D CNNs, SwinUNETR, and hybrid transformer-CNN models achieve strong Dice or Jaccard scores. The second direction is abnormality classification, where self-supervised, contrastive, and transfer-learning methods help under limited annotations. The proposed project is positioned between these two directions by using segmentation as an anatomical guide for multi-label abnormality prediction.

This state-of-the-art framing fits the project because the system does not claim to replace radiological interpretation. Instead, it combines a validated segmentation-style workflow with localized ROI classification. This makes the model more interpretable than a full-volume classifier and more clinically relevant than a segmentation-only system.




Figure 3.1: Research gap and positioning of the proposed system




Figure 3.2: Literature-based justification of ConvNeXt-Tiny and Swin-T comparison

Table 3.1: Summary of related work in sinus CT analysis, backbones, and medical imaging

Paper	Approach	Architecture	Main	Strength	Limitation
Zhang et al., 2024, NasalSeg	Open-access CT dataset and benchmark for nasal cavity and paranasal sinus segmentation.	Dataset with voxel-level annotations; benchmark segmentation models.	Provides 130 3D CT scans with five annotated structures.	Public dataset directly related to this project.	Labels only five structures and do not directly include septum or turbinates.
	Transformer-	Swin		Strong	Memory

Hatamizadeh et al., 2022, SwinUNETR	based 3D medical image segmentation.	Transformer encoder with U-shaped decoder.	Captures long-range context in 3D segmentation.	architecture for volumetric medical segmentation.	intensive and originally validated on other medical tasks.
Liu et al., 2022, ConvNeXt	Modernized convolutional architecture.	Pure ConvNet with transformer-era design improvements.	Shows that modern CNNs can compete with transformer backbones.	Supports choosing ConvNeXt-Tiny for limited-data transfer learning.	Designed for natural images; medical transfer still requires careful ROI design.
Liu et al., 2021, Swin Transformer	Hierarchical transformer with shifted windows.	Window-based self-attention and multi-scale features.	Efficient transformer backbone for vision tasks.	Good comparison model against ConvNeXt.	Can require more data and careful regularization.
Whangbo et al., 2024	Multi-class paranasal sinus segmentation in sinusitis CT.	3D U-Net variants.	Segments major sinus regions in CT images.	Clinically relevant paranasal sinus segmentation setting.	Focuses on segmentation rather than abnormality prediction.
Bhattacharya et al., 2024	Self-supervised learning for paranasal anomaly classification.	3D convolutional autoencoder / self-supervised pipeline.	Classifies normal versus anomalous maxillary sinus patterns.	Addresses limited labeled data through self-supervision.	Narrower classification scope than multi-label abnormality prediction.
Ozturk et al., 2024	Automatic maxillary sinus segmentation on CBCT.	U-Net segmentation model.	Segments maxillary sinus structures.	Relevant to maxillary anatomy and dental/sinus imaging.	CBCT-focused and not a full multi-label CT abnormality pipeline.
Song et al., 2025	Comparison of CNNs, ViTs, and hybrid networks for sinus segmentation.	CNN, ViT, and hybrid segmentation networks.	Compares architecture families for paranasal sinus segmentation.	Useful for understanding architecture trade-offs.	Segmentation-focused and does not solve clinical multi-label classification.

Xia et al., 2025	Medical image classification using an improved ConvNeXt-Tiny.	Lightweight ConvNeXt-Tiny with pooling and attention refinements.	Shows ConvNeXt-Tiny can be adapted to medical classification.	Supports ConvNeXt-Tiny as a reasonable classification backbone.	Not specific to sinus CT or multi-label ROI classification.
Roy et al., 2023, MedNeXt	Medical image architecture based on ConvNeXt-style blocks.	3D ConvNeXt encoder-decoder.	Adapts ConvNeXt principles to medical imaging.	Supports modern CNNs in data-scarce medical settings.	Segmentation-focused, not a direct classifier for sinus CT.
Gulsoy and Kablan, 2025, FocalNeXt	CT lung cancer classification using ConvNeXt-augmented architecture.	ConvNeXt + FocalNet hybrid.	Applies ConvNeXt-related features to CT classification.	Shows ConvNeXt components are useful in CT classification.	Different disease and dataset.
Aburass et al., 2025	Review of vision transformers in medical imaging.	Review of ViT/Swin-style methods.	Summarizes transformer applications in medical classification and segmentation.	Justifies including Swin-T as a comparison.	Review-level evidence; not a sinus-specific experiment.
Wang et al., 2024	2.5D deep learning and conventional features for CT-based prediction.	2.5D masks/features plus model fusion.	Uses 2.5D representation to balance context and efficiency.	Supports 2.5D design for limited data.	Different clinical target and not sinus CT.

3.18 Comparative Analysis of Existing Systems

Table 3.2: Comparison of existing approaches and the proposed system

System / Approach	Features	Technologies	Strengths	Weaknesses
-------------------	----------	--------------	-----------	------------

Manual CT interpretation	Radiologist reviews full CT volume and reports findings.	Clinical expertise and PACS viewer.	Gold standard for clinical decision making.	Time-consuming and subjective variability can occur.
Segmentation-only AI	Automatically segments anatomical structures.	U-Net, nnU-Net, SwinUNETR, 3D CNNs.	Provides localization and quantitative anatomy.	Does not directly predict abnormality labels.
Image-level classification AI	Predicts disease label from full images or slices.	CNNs, ViTs, transfer learning.	Simple inference pipeline.	May ignore anatomy and learn irrelevant correlations.
Backbone-only transfer learning	Uses pretrained 2D networks on slices or projections.	ConvNeXt, Swin, EfficientNet, ResNet.	Efficient and strong baseline.	Needs careful input design to avoid losing anatomical context.
Proposed system	Segments anatomy, extracts ROIs, predicts five abnormalities.	SwinUNETR + ROI extraction + ConvNeXt-Tiny.	Combines localization and classification.	Limited by small dataset, indirect masks for septum/turbinates, and label imbalance.

3.19 Research Gap and Proposed Innovation

The reviewed work shows that strong methods exist for medical segmentation and image classification, but there remains a gap in segmentation-assisted multi-label abnormality prediction for sinus CT using the available NasalSeg anatomical structures. Existing sinus segmentation papers localize anatomy well, while paranasal classification papers usually focus on narrower maxillary anomaly tasks. The current project addresses this gap by combining SwinUNETR-based anatomical localization with ROI-based five-label abnormality classification.

The proposed innovation is to use the available anatomical masks as stable context regions. The final classifier does not attempt to create unreliable artificial masks for unsegmented structures. Instead, it uses clinically motivated neighboring ROIs: the septum is modeled through the region between both nasal cavities, and inferior turbinate hypertrophy is modeled through the ipsilateral nasal cavity ROI.

The classification-backbone literature supports comparing ConvNeXt-Tiny with Swin-T, but the experimental interpretation is that Swin-T is a strong transformer comparison rather than the final model. ConvNeXt-Tiny is selected as the final classifier because it achieved stronger

performance in the project experiments while maintaining a suitable balance between transfer learning, stability, and limited-data behavior.

3.20 Chapter Summary

This chapter reviewed the main literature related to sinus CT segmentation, paranasal anomaly classification, modern CNN and transformer backbones, and 2.5D medical imaging. The review shows that the proposed project is positioned between segmentation-only sinus CT systems and image-level classification systems. Its main novelty is the use of segmentation-guided ROIs to support multi-label abnormality prediction while explicitly acknowledging the limitations caused by missing direct septum and turbinate masks.

Chapter Four: Research Methodology

4.1 Introduction

This chapter describes the methodology used to develop the AI system for sinus CT analysis. The pipeline begins with CT volume preprocessing and anatomical segmentation. The segmentation output is used to extract fixed-size regions of interest. These ROIs are converted into 2.5D views and passed to a multi-label classifier. The final system predicts five abnormality labels and selects ConvNeXt-Tiny as the final classification model.




Figure 4.1: Overall pipeline of the proposed sinus CT analysis system

4.2 Dataset Description and Annotation Characteristics

The project uses CT volumes and segmentation labels from the NasalSeg dataset together with the provided clinical annotation workbook. The segmentation task uses 130 CT cases with five annotated anatomical structures: right maxillary sinus, left maxillary sinus, right nasal cavity, left nasal cavity, and nasopharynx. These anatomical labels were used for localization and ROI extraction, not as direct abnormality labels.

For classification, the annotation workbook was converted into a model-ready multi-label table. Each selected clinical finding was mapped to an independent binary target, because one CT case may contain several abnormalities at the same time. The final diagnostic targets were nasal

septum abnormality, right maxillary abnormality, left maxillary abnormality, right inferior turbinate hypertrophy, and left inferior turbinate hypertrophy.

The annotation workbook originally contained more detailed textual or categorical descriptions for each anatomical region. However, these detailed values **were not used** as separate classes because the dataset size was limited and the positive cases within each region were already small. For example, even if an anatomical region contained around 60 positive cases, dividing these cases into many abnormality subtypes would create extremely small and imbalanced subclasses. This would make supervised learning unstable and would also make evaluation unreliable.

Therefore, the diagnostic annotations were reformulated into binary present/absent labels for each selected anatomical target. This design preserved the clinically relevant question of whether an abnormality exists in each region while avoiding excessive class fragmentation.

The Normal column was not used as a separate sixth output. Instead, a case without any of the selected abnormality labels was represented by an all-zero target vector. This avoided mixing a normal/not-normal decision with the five abnormality-specific outputs and kept the learning problem consistent with multi-label classification.

The workbook also revealed a strong imbalance between labels. Nasal septum abnormality and inferior turbinate hypertrophy had substantially more positive cases than the maxillary abnormality labels, which contained only 22 right maxillary and 23 left maxillary positives. This imbalance influenced the later choice of focal loss, positive class weighting, threshold tuning, and 5-fold out-of-fold evaluation.

A second important characteristic was the mismatch between available segmentation masks and some diagnostic targets. Maxillary labels had direct anatomical masks, but nasal septum and inferior turbinate labels did not. Therefore, the final classifier used contextual ROIs: both nasal cavities for septum-related prediction and the ipsilateral nasal cavity for inferior turbinate hypertrophy.

Table 4.1: Dataset summary

Item	Value
Total matched CT cases	130
Normal cases in annotation workbook	18

Cases with at least one selected abnormal label	107
Final classification task	Five-label multi-label abnormality prediction
Segmentation classes	Background + five anatomical structures
Classification evaluation	Balanced 5-fold out-of-fold evaluation

Table 4.2: Final target labels and positive sample counts

Target label	Positive count	Interpretation
nasal_septum_abnormal	71	Any non-empty septum abnormality annotation mapped to positive.
right_maxillary_abnormal	22	Right maxillary sinus abnormality present.
left_maxillary_abnormal	23	Left maxillary sinus abnormality present.
right_inferior_turbinate_hypertrophy	60	Right inferior turbinate hypertrophy present.
left_inferior_turbinate_hypertrophy	60	Left inferior turbinate hypertrophy present.




Figure 4.2: Label preparation workflow from the annotation sheet

4.3 Data Preprocessing

For segmentation, CT volumes are loaded as 3D medical images and transformed using MONAI. The main preprocessing operations include orientation handling, spacing normalization, intensity windowing, scaling, and conversion to tensors. The HU window used in the segmentation notebook is -400 to 1000, which keeps relevant sinonasal soft tissue and bony structures while limiting extreme intensity values.

For classification, each CT and segmentation pair is used to extract anatomy-guided ROIs. Each ROI includes two channels: the CT intensity crop and the corresponding segmentation or mask context. The ROI crop is resized to a fixed size, then converted into 2.5D images by taking representative axial, coronal, and sagittal planes. These three planes form a three-channel image suitable for pretrained 2D backbones.

4.4 Data Augmentation

Data augmentation was applied only to the training data. For segmentation, augmentation included random crops, flips, rotations, Gaussian noise, and intensity shift or scale transformations. For classification, augmentation was applied to improve robustness of the 2.5D ROI views without changing the target labels. Validation and test samples were kept stable except for deterministic preprocessing so that evaluation remained unbiased.

4.5 Data Partitioning and 5-Fold Evaluation

The segmentation notebook used an executed split of 91 training cases, 19 validation cases, and 20 test cases. The classification notebook used balanced 5-fold out-of-fold evaluation across the 130 matched cases. This setup ensured that each case was evaluated by a model that did not train on it. The 5-fold design was selected because the dataset is limited and a single train/test split could produce unstable results, especially for rare labels such as maxillary abnormalities.

Table 4.3: Data partitioning strategy used in this project

Stage	Partitioning	Purpose
Segmentation	91 train / 19 validation / 20 test	Train SwinUNETR and evaluate anatomical segmentation on held-out cases.
Classification	Balanced 5-fold OOF	Compare classifiers while using every case once as validation.
Threshold tuning	OOF predictions	Select label-specific probability thresholds for final metrics.

4.6 SwinUNETR Segmentation Stage

The segmentation stage uses SwinUNETR, a 3D transformer-based U-shaped architecture. The model is trained to predict six classes: background, right maxillary sinus, left maxillary sinus, right nasal cavity, left nasal cavity, and nasopharynx. The training objective is DiceCE loss, which combines overlap-based Dice loss and cross-entropy loss. This is suitable for segmentation because it encourages both voxel-level class prediction and mask overlap.

Table 4.4: SwinUNETR segmentation training configuration

Parameter	Value
Architecture	SwinUNETR
Framework	MONAI / PyTorch
Input spacing	$1.0 \times 1.0 \times 1.5$
ROI size	$96 \times 96 \times 32$
HU window	-400 to 1000
Loss function	DiceCE loss
Optimizer	AdamW
Validation method	Sliding-window inference
Executed split	91 train / 19 validation / 20 test cases




Figure 4.3: Augmentation sanity-check visualization

4.7 ROI Extraction Strategy

The final classification model uses five stable contextual ROIs. The right and left maxillary labels use the corresponding maxillary sinus masks directly. The nasal septum label uses a contextual ROI derived from both nasal cavities with padding. This design follows the anatomical fact that the nasal septum lies between the two nasal cavities, while avoiding brittle pseudo-septum masks. The inferior turbinate labels use the ipsilateral nasal cavity ROIs because direct turbinate masks are not available in the dataset.

Table 4.5: ROI-to-label mapping used in the final classification system

Target label	ROI name	Segmentation labels used	Reasoning
right_maxillary_abnormal	right_maxillary_roi	Right maxillary sinus	Direct anatomical match.
left_maxillary_abnormal	left_maxillary_roi	Left maxillary sinus	Direct anatomical match.
nasal_septum_abnormal	nasal_septum_region_roi	Right + left nasal cavities	Septum lies between both nasal cavities; contextual ROI is more stable.
right_inferior_turbinate_abnormal	right_nasal_cavity_roi	Right nasal cavity	Inferior turbinate is located inside the nasal cavity; no direct turbinate mask available.
left_inferior_turbinate_abnormal	left_nasal_cavity_roi	Left nasal cavity	Inferior turbinate is located inside the nasal cavity; no direct turbinate mask available.




Figure 4.4: Final ROI definitions used in the classification stage




Figure 4.5: Example ROI extraction results from the classification notebook

4.8 2.5D ROI Generation

Each 3D ROI was converted into a 2.5D representation by taking representative axial, coronal, and sagittal views. These three views were stacked as a three-channel image. This design preserved partial volumetric context while allowing the use of pretrained 2D image classification backbones. It also reduced the memory cost compared with fully 3D classifiers, which is important for a dataset of 130 cases.




Figure 4.6: 2.5D ROI generation from axial, coronal, and sagittal views

4.9 Important Exploratory Experiments

The classification stage went through several major experimental phases before reaching the current ConvNeXt-Tiny result. This section documents the important development path because the final system was not obtained from a single direct experiment. The project first established anatomical segmentation, then tested ROI-based 3D classification, then moved toward 2.5D transfer learning after several limitations became clear.

The first foundation was the SwinUNETR segmentation model. It segmented the available NasalSeg anatomical structures and provided the masks used to crop regions of interest. This changed the classification task from blind whole-volume prediction into anatomy-guided prediction, where each label could be linked to a clinically meaningful region.

After segmentation, the diagnosis problem was reformulated as multi-label classification. This was necessary because a single sinus CT case can contain septal abnormality, turbinate hypertrophy, and maxillary sinus abnormality at the same time. The model therefore used independent sigmoid outputs rather than a single softmax class.

The first diagnostic models used segmentation-guided 3D ROI crops with a ResNet-style classifier. Early experiments included seven abnormality labels, but the middle turbinate labels had too few positive cases and produced unstable zero or near-zero F1 scores. Therefore, the target set was reduced to the five main labels used in the final system.

The 3D ROI classifier was then improved using two-channel inputs. Each ROI contained both the CT intensity crop and the corresponding mask crop, allowing the model to learn image appearance and anatomical location together. Separate label-specific heads were also more stable than shared bilateral heads, because shared turbinate heads tended to increase false positives.

Several imbalance and augmentation strategies were tested. Raw positive class weighting was too aggressive, so softened square-root positive weights and capped weights were explored. Light on-the-fly CT intensity augmentation was useful, while saved augmented ROI copies were rejected because they mostly created near-duplicate versions of the same patients rather than true new clinical diversity.

Threshold tuning also became necessary. A fixed threshold of 0.5 was not reliable for all labels because each abnormality had a different prevalence and probability distribution. Hybrid and validation-based thresholds improved some fixed-split results, but the results remained sensitive to the split, which motivated the move toward 5-fold out-of-fold evaluation.

The best fixed-split ROI-based 3D ResNet model achieved an approximate macro F1 around 0.63, but later cross-validation showed that this was likely optimistic. In the more honest 5-fold evaluation, the ROI 3D ResNet approach reached an out-of-fold macro F1 of approximately 0.588, confirming that the custom 3D classifier had moderate capability but also a clear performance ceiling.

Additional diagnostic experiments were useful even when they were not selected. Handcrafted ROI statistics, tighter pseudo-septum crops, bilateral context, shared heads, signed and absolute difference features, and ground-truth mask extraction were all examined. These experiments showed that segmentation quality was helpful but not the only bottleneck; the main limitations were small sample size, rare positive labels, indirect septum and turbinate targets, and threshold instability.

The major improvement came from replacing the custom 3D encoder with a 2.5D transfer-learning classifier. Each 3D ROI was converted into axial, coronal, and sagittal views, enabling the use of pretrained image backbones. ConvNeXt-Tiny was the strongest direction because it improved feature extraction while remaining suitable for the limited dataset size.

In the later focused experiments, ConvNeXt-Tiny with focal loss gamma equal to 1, positive class weighting, light augmentation, strict early stopping, ReduceLROnPlateau, and threshold tuning achieved the strongest fold-level macro F1. Swin-T was retained as an important transformer comparison experiment, but it did not outperform the final ConvNeXt-Tiny configuration.

Overall, the experimental history shows that the final pipeline was shaped by systematic failure analysis as much as by successful results. The project moved from segmentation, to ROI-based 3D diagnosis, to honest 5-fold evaluation, and finally to 2.5D ConvNeXt-Tiny transfer learning because each stage exposed a specific limitation of the previous design.

Timeline of Major Classification Experiments




Figure 4.7: Timeline of major classification experiments

Table 4.6: Timeline of major exploratory experiments and outcomes

Stage	Experiment / Idea	Main observation	Decision
1	SwinUNETR segmentation foundation	Mean Dice about 0.87; masks localized maxillary sinuses, nasal cavities, and nasopharynx.	Kept as anatomical localization stage.
2	Annotation cleaning and multi-label formulation	Workbook findings were converted into five independent abnormality labels; normality was represented by all-zero targets.	Adopted for all diagnostic experiments.
3	Initial seven-label ROI 3D ResNet	First full segmentation-guided diagnostic classifier; middle turbinate labels failed because of very small positive counts.	Reduced task to five labels.
4	Five-label 3D ROI classifier with CT + mask input	Two-channel ROI input and separate label heads gave the most stable 3D ResNet direction.	Kept as main 3D baseline.
5	Imbalance handling and augmentation	Softened positive weights, light CT augmentation, and label-specific thresholds reduced some instability.	Retained only stable components.
6	Best fixed-split 3D ROI	Approximate macro F1 around 0.63, but	Used as fixed-split

	ResNet	performance depended strongly on the selected split and thresholds.	baseline only.
7	Five-fold 3D ROI ResNet evaluation	OOF macro F1 about 0.588; showed that fixed-split results were optimistic.	Adopted 5-fold OOF evaluation.
8	Alternative ROI and head designs	Handcrafted features, tighter septum ROIs, bilateral context, shared heads, and difference features did not consistently improve results.	Rejected in favor of simpler stable design.
9	Ground-truth mask experiment	Oracle masks improved some ranking behavior but did not clearly improve thresholded OOF F1.	Confirmed labels/data were main bottleneck.
10	2.5D transfer-learning transition	Converted ROIs into axial, coronal, and sagittal views and used pretrained backbones.	Adopted for final classifiers.
11	ConvNeXt-Tiny focal-loss experiments	ConvNeXt-Tiny with focal loss $\gamma=1$ produced the strongest fold-level macro F1 among tested backbones.	Selected as final direction.
12	Swin-T transformer comparison	Swin-T was competitive and useful as a transformer baseline, but remained below ConvNeXt-Tiny.	Kept as comparison, not final model.




Figure 4.8: Macro F1 comparison across experiments.

4.10 Final Selected Classification Model

The final selected classification model is ConvNeXt-Tiny. Swin-T was retained as an important transformer-based comparison during the exploratory and model-selection stages, but it is not presented as a second final model. The final ConvNeXt-Tiny classifier uses the same five binary abnormality targets, segmentation-guided ROI cache, 2.5D axial-coronal-sagittal representation, focal loss with gamma equal to 1, positive class weighting, light augmentation, strict early stopping, and 5-fold out-of-fold evaluation.

4.11 Implementation Environment

The experiments were implemented in a notebook-based Python environment. The segmentation pipeline used PyTorch and MONAI for 3D medical image processing, transforms, training, and sliding-window inference. The classification pipeline used PyTorch, TorchVision, pandas,

NumPy, scikit-learn, and Matplotlib. Training was performed in a Kaggle GPU environment. The final notebooks produced checkpoint files, ROI caches, OOF prediction tables, threshold tables, metrics CSV files, and plots used for analysis.

4.12 Chapter Summary and Research Contribution

The methodology combines 3D anatomical segmentation with ROI-based 2.5D multi-label classification. The major methodological decision is to avoid using the full CT volume as an unlocalized classifier input and instead use segmentation-guided anatomical crops. This makes the pipeline more interpretable and allows each output label to be associated with a clinically meaningful region. The final selected classifier is ConvNeXt-Tiny, while Swin-T is documented as a previous transformer comparison experiment rather than a second final model.

Chapter Five: Research Results

5.1 Introduction

This chapter presents the experimental results of the proposed sinus CT analysis system. The results are reported in two main parts: segmentation results and classification results. Segmentation results evaluate the SwinUNETR model as the anatomical localization stage. Classification results evaluate the final multi-label abnormality prediction stage and compare ConvNeXt-Tiny with Swin-T.

5.2 Tools and Evaluation Methods

The experiments were implemented using Python, PyTorch, MONAI, Torchvision, NumPy, pandas, scikit-learn, and Matplotlib. Training was performed in a Kaggle GPU environment. The segmentation model was evaluated with Dice, IoU, precision, recall, inference time, and throughput. The classification models were evaluated with macro accuracy, macro precision, macro recall, macro specificity, macro balanced accuracy, macro F1, macro AUC, macro average precision, overall label accuracy, and exact match accuracy.

5.2.1 Experimental Environment

The segmentation and classification experiments were developed as notebook-based workflows. This allowed step-by-step inspection of data loading, preprocessing, ROI extraction, model training, metric computation, and figure generation. The final artifacts included trained checkpoints, ROI cache files, fold summaries, out-of-fold predictions, tuned thresholds, per-label metrics, and plots.

5.2.2 Evaluation Metrics Used in the Final Experiments

Segmentation was evaluated using mean Dice, mean IoU, precision, recall, inference time, and throughput. Classification was evaluated using label-wise and macro metrics. Because the task is multi-label and imbalanced, macro F1 and out-of-fold metrics were emphasized over simple accuracy. Exact match accuracy was also reported, but it is naturally strict because it requires all five labels of each case to be correct at the same time.

5.3 Segmentation Results

The SwinUNETR model achieved strong segmentation performance on the test set. The mean Dice score was 0.8743 and the mean IoU was 0.8102. The high recall value indicates that the model generally captured most of the reference anatomical regions. The segmentation results are suitable for the intended purpose of guiding ROI extraction in the classification stage.

5.3.1 Quantitative Segmentation Results




Figure 5.1: SwinUNETR training and validation curves

Table 5.1: Overall segmentation performance on the test set

Metric	Value
Mean Dice	0.8743
Mean IoU	0.8102
Mean Precision	0.8424
Mean Recall	0.9158
Inference time	0.34 s / volume
Throughput	2.90 volumes / s

Table 5.2: Per-class segmentation performance

Class	Dice	IoU	Precision	Recall
Background	0.9896	0.9796	0.9946	0.9846
Right maxillary sinus	0.9107	0.8543	0.8797	0.9476
Left maxillary sinus	0.8697	0.8104	0.8503	0.9053
Right nasal cavity	0.8009	0.7010	0.7440	0.8729
Left nasal cavity	0.8079	0.7138	0.7520	0.8772
Nasopharynx	0.8669	0.8024	0.8338	0.9075

5.3.2 Qualitative Segmentation Results

Qualitative examples should be used to confirm that the segmentation mask follows the expected anatomical structures and that failure cases are not hidden by global metrics. These images will be inserted later from the notebook outputs or manually prepared visualizations.




Figure 5.2: Example segmentation prediction on a test case

5.4 Classification Results

The final classification evaluation focuses on ConvNeXt-Tiny as the selected model. Swin-T results are used only as supporting evidence from the model-selection process, showing why the convolutional ConvNeXt-Tiny backbone was selected over the transformer comparison model for this limited ROI-based dataset. ConvNeXt-Tiny achieved the strongest reported performance, with a fold-level macro F1 of 0.7849 and an out-of-fold macro F1 of 0.6928.

Table 5.3: Supporting model-selection comparison between ConvNeXt-Tiny and Swin-T

Model	Fold-level macro F1	OOF macro F1	Overall label accuracy	Exact match accuracy
ConvNeXt-Tiny	0.7849	0.6928	0.7877	0.3846
Swin-T	0.7401	0.6565	0.7031	0.1923

5.4.1 Final ConvNeXt-Tiny Results

ConvNeXt-Tiny achieved the best overall classification performance. Its OOF macro F1 of 0.6928 is the main final score because every case is evaluated by a model that did not train on it. Its overall label accuracy was 0.7877, and exact match accuracy was 0.3846. Exact match is strict in this multi-label task because a case with four correct labels and one wrong label is still considered incorrect under this metric.




Figure 5.3: ConvNeXt-Tiny loss curves




Figure 5.4: ConvNeXt-Tiny macro F1 curves

Table 5.4: OOF per-label metrics for ConvNeXt-Tiny

Label	Threshold	Accuracy	Precision	Recall	F1
left_inferior_turbinate_abnormal	0.49	0.7308	0.7049	0.7167	0.7107
left_maxillary_abnormal	0.42	0.8462	0.5600	0.6087	0.5833
nasal_septum_abnormal	0.44	0.6846	0.6786	0.8028	0.7355
right_inferior_turbinate_abnormal	0.46	0.7769	0.7385	0.8000	0.7680
right_maxillary_abnormal	0.43	0.9000	0.7647	0.5909	0.6667

5.4.2 Per-label Performance Analysis

The per-label results show that the model performed better on some anatomical targets than others. Right inferior turbinate hypertrophy and nasal septum abnormality achieved relatively strong F1 scores. The maxillary labels were less stable, especially left maxillary abnormality. This does not mean that the ROI design was incorrect; rather, it reflects the small positive counts and frequent co-occurrence of maxillary abnormalities with septal or turbinate findings.

5.4.3 Confusion Matrix and Multi-Label Interpretation

The final notebook includes a 5×5 dominant-label confusion matrix for ConvNeXt-Tiny. These matrices are useful for visualizing the dominant predicted abnormality among abnormal cases, but they must be interpreted carefully. The classification task is multi-label, while the matrix forces each abnormal case into a single dominant true label and a single dominant predicted label. Therefore, a case can be partially correct in the multi-label sense but still appear as a confusion in the dominant-label matrix.

This issue is especially important for maxillary abnormalities. Maxillary-positive cases are relatively rare and often co-occur with septal or turbinate abnormalities. As a result, if the model correctly predicts a maxillary label but assigns higher confidence to a co-occurring septal or turbinate label, the dominant-label matrix may underrepresent the true multi-label correctness. For this reason, the report emphasizes macro F1, per-label metrics, and OOF evaluation in addition to the 5×5 visualization.




Figure 5.5: ConvNeXt-Tiny 5×5 dominant-label confusion matrix

5.4.4 ROC and Precision-Recall Analysis

ROC and precision-recall curves provide additional insight into threshold-dependent behavior. ROC curves show how sensitivity and specificity trade off across probability thresholds. Precision-recall curves are especially important for imbalanced labels because they focus on positive-class retrieval. ConvNeXt-Tiny showed useful separability for several labels, although maxillary labels remained harder because of lower positive counts.




Figure 5.6: ROC curves for the final ConvNeXt-Tiny model




Figure 5.7: Precision-Recall curves for the final ConvNeXt-Tiny model

5.5 Comparison with Previous Works

The proposed system differs from many previous works because it combines anatomical segmentation with multi-label abnormality prediction. NasalSeg and similar studies focus mainly on segmentation, while several maxillary-sinus studies focus on a narrower classification target. This project uses segmentation as an intermediate anatomical guide and then predicts

multiple abnormality labels. This is a more difficult setting because labels may co-occur and some targets do not have direct segmentation masks.

Table 5.5: Comparison of the proposed model with previous works

Work	Task	Data / Scope	Main Result or Focus	Relation to this project
NasalSeg	Anatomical segmentation	130 CT scans, five structures	Public dataset and segmentation benchmark	Provides the anatomical basis for this project.
Whangbo et al., 2024	Paranasal sinus segmentation	Sinusitis CT data	Multi-class sinus region segmentation	Related segmentation task but not multi-label classification.
Bhattacharya et al., 2024	Anomaly classification	Paranasal/maxillary anomaly focus	Self-supervised representation learning	Related classification direction but narrower label scope.
ConvNeXt	General image classification backbone	Natural image benchmarks	Modern CNN representation learning	Used as final classifier backbone.
Swin Transformer	General transformer vision backbone	Natural image / dense prediction benchmarks	Hierarchical shifted-window attention	Used as final transformer comparison.
Improved ConvNeXt-Tiny medical classification	Medical image classification	General medical image classification	Efficient ConvNeXt-Tiny adaptation	Supports ConvNeXt-Tiny choice but not sinus-specific.
Proposed system	Segmentation-guided multi-label classification	130 CT cases and five abnormality labels	ConvNeXt-Tiny OOF macro F1 = 0.6928	Combines anatomical localization and abnormality prediction.

5.6 Discussion of Results

5.6.1 Effect of Segmentation-Guided ROI Extraction

Segmentation-guided ROI extraction was one of the main reasons the final pipeline became interpretable. Instead of asking the classifier to search the whole CT volume, each label was connected to a specific region. This design is especially important in sinus CT because abnormalities are region-dependent. The segmentation stage therefore acts as an anatomical filter

before classification.

5.6.2 Effect of Final ROI Design

The final ROI design avoided fragile artificial masks for structures that NasalSeg does not directly annotate. The septum was modeled using both nasal cavities with padding, while inferior turbinate hypertrophy was modeled using the ipsilateral nasal cavity. These contextual ROIs did not eliminate all difficulty, but they were more stable than pseudo-septum or pseudo-turbinate strategies.

5.6.3 Effect of Imbalance Handling

Class imbalance was handled with focal loss and positive class weighting. These choices helped the model pay more attention to difficult and underrepresented labels. However, imbalance was not fully solved, especially for maxillary abnormalities. This shows that algorithmic weighting can reduce imbalance effects but cannot replace the need for more positive cases.

5.6.4 Why ConvNeXt-Tiny Was Selected

During model selection, ConvNeXt-Tiny outperformed the Swin-T comparison in fold-level macro F1, out-of-fold macro F1, overall label accuracy, and exact match accuracy. This result suggests that a modern convolutional backbone was better suited to the available localized ROI data than the transformer comparison model. Therefore, ConvNeXt-Tiny is selected as the final proposed classifier.

The literature review supports this interpretation. ConvNeXt was designed to modernize CNNs while preserving convolutional inductive bias, and several medical-imaging works have adapted ConvNeXt-style designs to limited-data medical settings. Therefore, the choice of ConvNeXt-Tiny is supported both experimentally in this project and conceptually by recent medical-imaging research.

5.6.5 Maxillary Label Instability

The right and left maxillary labels were the most unstable labels. The main reason is the low number of positive cases: 22 for right maxillary abnormality and 23 for left maxillary abnormality. These labels also frequently co-occur with septal or turbinate abnormalities. Co-

occurrence makes it harder to isolate the contribution of each label and can make the dominant-label confusion matrix appear worse than the full multi-label prediction actually is.

5.6.6 Limitations of the 5×5 Dominant-Label Confusion Matrix

The 5×5 matrix is useful as a compact visualization, but it is not the main selection metric. It compresses a multi-label case into one dominant true label and one dominant predicted label. This can hide partially correct predictions. For example, a case with both maxillary and septal abnormalities may be counted as confused if the model predicts the correct maxillary label but gives higher confidence to the co-occurring septal label. Therefore, OOF macro F1, per-label F1, and threshold-tuned metrics are more appropriate for evaluating the final model.

5.7 Chapter Summary

The results show that SwinUNETR provides strong anatomical segmentation and that ConvNeXt-Tiny is the best final classifier among the evaluated models. The main remaining weaknesses are related to label imbalance, multi-label co-occurrence, and lack of direct masks for septum and turbinates. These issues guide the future work described in the final chapter.

Chapter Six: Conclusion

6.1 Introduction

This chapter summarizes the main findings of the project, states the conclusions, identifies the research limitations, and proposes future directions. The project should be viewed as a research prototype for segmentation-guided sinus CT analysis rather than a clinically deployed system.

6.2 Conclusions

This project developed a segmentation-assisted AI system for sinus CT analysis. The system uses SwinUNETR to segment key anatomical structures, extracts anatomy-guided regions of interest, and applies multi-label classification to predict five selected abnormalities. ConvNeXt-Tiny was selected as the final classifier because it achieved the best overall performance among the final compared models.

The segmentation model achieved strong performance, with a mean Dice score of 0.8743 and mean IoU of 0.8102 on the test set. These results indicate that the model can localize the available NasalSeg structures well enough to support downstream ROI extraction. The classification model achieved a ConvNeXt-Tiny OOF macro F1 score of 0.6928, outperforming Swin-T on the final evaluation. The results demonstrate that the proposed approach is promising but also data-limited.

The study also shows that evaluation must be interpreted carefully in multi-label medical imaging tasks. The 5×5 dominant-label confusion matrix is visually helpful but underrepresents multi-label correctness, especially for cases with co-occurring maxillary, septal, and turbinate abnormalities. Therefore, per-label and out-of-fold metrics are more appropriate for final model selection.

6.3 Research Limitations

- The dataset contains only 130 matched CT cases, which limits the generalization of deep learning models.

- The right and left maxillary abnormality labels have low positive counts compared with other labels.
- NasalSeg does not directly segment nasal septum or turbinate structures, requiring contextual ROI approximations.
- The classification task is multi-label, and co-occurrence makes dominant-label visualization incomplete.
- The final model was not externally validated on data from another hospital or scanner.
- The System is research prototype and should not be considered clinically deployable without further validation.

6.4 Challenges and Future Work

- Increase the dataset size, especially for right and left maxillary abnormality labels.
- Add direct expert annotations for the nasal septum and turbinates to reduce dependence on contextual proxy ROIs.
- Validate the final model on an external dataset from another scanner or hospital.
- Explore 3D or hybrid 2.5D/3D classifiers after increasing the number of labeled cases.
- Improve label quality through radiologist review and standardized abnormality definitions.
- Evaluate calibration and uncertainty estimation to make predictions more clinically interpretable.
- Develop a user interface that displays CT slices, segmentation masks, ROIs, and predicted abnormalities together.
- Use multi-label evaluation dashboards rather than relying only on dominant-label confusion matrices.

Table 6.1: Summary of main findings and future work directions

Finding	Implication	Future direction
SwinUNETR segmentation performed well	Segmentation can guide ROI extraction	Use improved/manual masks for more structures.
ConvNeXt-Tiny outperformed Swin-T	Modern CNNs were more stable for this dataset	Keep ConvNeXt-Tiny as final baseline.

Maxillary labels were unstable	Low positive counts and co-occurrence affect learning	Collect more maxillary-positive cases.
Dominant-label matrix is limited	Multi-label correctness may be hidden	Use per-label and case-level multi-label metrics.
Septum/turbinate masks were unavailable	Contextual ROIs are approximations	Annotate septum and turbinate structures directly.

References

- [1] Y. Zhang et al., “NasalSeg: A Dataset for Automatic Segmentation of Nasal Cavity and Paranasal Sinuses from 3D CT Images,” Scientific Data, vol. 11, 1329, 2024. <https://doi.org/10.1038/s41597-024-04176-1>
- [2] Y. Zhang, “NasalSeg Dataset for Nasal Cavity and Paranasal Sinuses Segmentation from CT Images,” Zenodo, 2024. <https://doi.org/10.5281/zenodo.12177181>
- [3] A. Hatamizadeh et al., “Swin UNETR: Swin Transformers for Semantic Segmentation of Brain Tumors in MRI Images,” Brainlesion, 2022. https://doi.org/10.1007/978-3-031-08999-2_22
- [4] Z. Liu et al., “A ConvNet for the 2020s,” Proceedings of CVPR, 2022. <https://arxiv.org/abs/2201.03545>
- [5] Z. Liu et al., “Swin Transformer: Hierarchical Vision Transformer Using Shifted Windows,” Proceedings of ICCV, 2021. <https://arxiv.org/abs/2103.14030>
- [6] J. Whangbo et al., “Deep Learning-Based Multi-Class Segmentation of the Paranasal Sinuses of Sinusitis Patients Based on Computed Tomographic Images,” Sensors, vol. 24, no. 6, 1933, 2024. <https://doi.org/10.3390/s24061933>
- [7] D. Bhattacharya et al., “Self-Supervised Learning for Classifying Paranasal Anomalies in the Maxillary Sinus,” International Journal of Computer Assisted Radiology and Surgery, 2024. <https://doi.org/10.1007/s11548-024-03172-5>
- [8] D. Bhattacharya et al., “Supervised Contrastive Learning to Classify Paranasal Anomalies in the Maxillary Sinus,” arXiv, 2022. <https://arxiv.org/abs/2209.01937>
- [9] B. Ozturk et al., “Automatic Segmentation of the Maxillary Sinus on Cone Beam Computed Tomographic Images with U-Net Deep Learning Model,” European Archives of Oto-Rhino-Laryngology, 2024. <https://doi.org/10.1007/s00405-024-08870-z>
- [10] D. Song et al., “Comparison of Segmentation Performance of CNNs, Vision Transformers, and Hybrid Networks for Paranasal Sinuses with Sinusitis on CT Images,” Scientific Reports, 2025. <https://doi.org/10.1038/s41598-025-17266-w>

-
- [11] J. Xia, Y. Yin, and X. Li, “An Efficient Medical Image Classification Method Based on a Lightweight Improved ConvNeXt-Tiny Architecture,” arXiv, 2025. <https://arxiv.org/abs/2508.11532>
- [12] S. Roy et al., “MedNeXt: Transformer-driven Scaling of ConvNets for Medical Image Segmentation,” MICCAI, 2023. https://doi.org/10.1007/978-3-031-43901-8_39
- [13] T. Gulsoy and E. B. Kablan, “FocalNeXt: A ConvNeXt augmented FocalNet architecture for lung cancer classification from CT-scan images,” Expert Systems with Applications, vol. 261, 125553, 2025. <https://doi.org/10.1016/j.eswa.2024.125553>
- [14] S. Aburass et al., “Vision Transformers in Medical Imaging: A Comprehensive Review of Advancements and Applications Across Multiple Diseases,” 2025. <https://pmc.ncbi.nlm.nih.gov/articles/PMC12701147/>
- [15] P. Wang et al., “Combining 2.5D deep learning and conventional features for early hematoma expansion prediction in spontaneous intracerebral hemorrhage,” Scientific Reports, 2024. <https://doi.org/10.1038/s41598-024-73415-7>
- [16] M. J. Cardoso et al., “MONAI: An open-source framework for deep learning in healthcare,” 2022. <https://monai.io/>
- [17] PyTorch Contributors, “PyTorch Documentation.” <https://pytorch.org/docs/stable/index.html>
- [18] scikit-learn Developers, “scikit-learn: Machine Learning in Python Documentation.” <https://scikit-learn.org/stable/>